

Causal Inference (6.S059/15.Co8/17.Co8)

Recitation, Week 7.

Topic: Supervised Machine Learning

Licheng Liu

April 5, 2024

MIT

Table of contents


1. Overview of Concepts

2. Supervised Learning Models and Algorithms

3. Practice in R and Python

1/ Overview of Concepts


Machine Learning and Causal Inference



Supervised Learning



- Algorithms that experience “features” and a supervision signal (target or label). **Labeled Set** of input-output pairs ($D = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$).
 - Using this data we build a **prediction** model or learner ($p(Y|\mathbf{X})$).
 - Good Learner: accurately predicts the target (*Performance Measure*).
 - Learner: Function $f(\mathbf{X})$ for predicting Y given the input vector \mathbf{X} . This requires a **Loss Function** ($L(Y, f(\mathbf{X}))$) for penalizing errors in prediction.
- **Type of Outcome:** Regression vs Classification.

Workflow



- **Generalization:** Algorithm must perform well on *new, previously unseen* inputs.
- Minimization of Training Error → Optimization Problem.
- Minimization of Generalization (Test) Error is the goal.

Model's Capacity, Bias-Variance Decomposition and Overfitting




- Ability to fit a wide variety of functions:
 - Underfitting.
 - Overfitting.
- Another Tradeoff: Flexibility (Prediction Accuracy) and Interpretability.


Other Relevant Concepts

- **Cross-Validation:** resampling method that uses different portions of the data to test and train a model on different iterations.

The technique of S -fold cross-validation, illustrated here for the case of $S = 4$, involves taking the available data and partitioning it into S groups (in the simplest case these are of equal size). Then $S - 1$ of the groups are used to train a set of models that are then evaluated on the remaining group. This procedure is then repeated for all S possible choices for the held-out group, indicated here by the red blocks, and the performance scores from the S runs are then averaged.



- **Curse of Dimensionality:** increasing data dimensions and its explosive tendencies. We would need an exponentially large quantity of training data (and computational power) to explore the high-dimensional feature space.




2/ Supervised Learning Models and Algorithms

Algorithms

- **No Free Lunch Theorem:** averaged over all possible DGPs, every algorithm has the same error rate when classifying previously unobserved points.
- No ML algorithm is universally any better than any other.
- Best performance: capacity is appropriate for the true complexity of the task and the amount of training data.

Algorithms



Regularization

- Any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.
- We can regularize a model that learns a function $f(x)$ by adding a penalty (regularizer) to the loss function (controls model complexity → overfitting).


$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s \quad (6.8)$$

and


$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s, \quad (6.9)$$

Regularization

- Ridge, LASSO, Elastic Net.
- LASSO for variable selection ($\beta_j^2 \rightarrow |\beta_j|$).



Non-Linear (more Flexible) Algorithms



3/ Practice in R and Python

Cross-Validation

- How to split data?
- Approach 1: randomly assign folds to each observation.
- This does not ensure that the size of each fold is equal.
- Practically ok if your N is large enough compared to K.

R Code

```
set.seed(02139)
# 5-folds cross validation with 10 observations
K <- 5; N <- 10; x <- rnorm(N); y <- rnorm(N)
folds <- sample(1:K, size = N, replace = TRUE) # assign folds
x[folds == 5]; y[folds == 5] # fifth folds: four observation
## [1] -0.07963674 -0.81726793 -0.16561194 0.36658112
## [1] 0.64319207 0.18791807 -0.05517906 0.63358473
x[folds == 2]; y[folds == 2] # second folds: one observations
## [1] -0.232987
## [1] 1.295322
```

Cross-Validation

- How to make the fold size equal?
- Approach 2: randomly assign IDs to each observation and select these observations.
- By doing so, you can ensure that the size is equal.

R Code

```
# assign folds
IDs <- sample(1:N, size = N, replace = FALSE)
# there should be 2 obs. in each fold
# ID 1:2 belongs to the first fold, 3:4 belongs to the second fold ...
x[IDs %in% c(1:2)] # first fold
## [1] 0.7720908 1.7165340
x[IDs %in% c(3:4)] # second fold
## [1] 0.9728744 0.3665811
```

Indicator function and Polynomials

- e.g., 4-th degree (order) polynomial regression

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \epsilon$$

- Terms x and x^2 will be highly correlated where $x > 0$
- `poly` function creates a “curved” set of variables from quadratic terms and makes sure that each term is not highly correlated to each other
- Technically, `poly` with the default setting computes “orthogonal polynomials” where each R Code polynomial is orthogonal

```
#Load data
data(cars)

#Run Model
m1 <- lm(dist ~ speed + I(speed^2), data = cars)
m2 <- lm(dist ~ poly(speed, degree = 2, raw = FALSE), data = cars)
m3 <- lm(dist ~ poly(speed, degree = 2, raw = TRUE), data = cars)

# Correlations
cor(model.matrix(m1)[, 2], model.matrix(m1)[, 3])
## [1] 0.9794765
cor(model.matrix(m2)[, 2], model.matrix(m2)[, 3])
## [1] 4.686464e-17
cor(model.matrix(m3)[, 2], model.matrix(m3)[, 3])
## [1] 0.9794765
```

LASSO

- Use the package `glmnet`.
 - `x`: input matrix.
 - `y`: target vector.
 - `family`: type of model (“binomial” for binary outcome).
 - `alpha`: mixing parameter (elastic net).
 - `lambda`: penalty parameter.
 - `standardize`: standardization of X (default = TRUE).
 - `exclude`: exclude variables from regularization/variable selection.
 - `maxit`: number of iterations.

R Code

```
#Load package
library(glmnet)

#Load data
data(QuickStartExample)
x <- QuickStartExample$x
y <- QuickStartExample$y

#Run learner
l1 <- glmnet(x=x, y=y, family = "gaussian", lambda = 0) #Ridge
l2 <- glmnet(x=x, y=y, family = "gaussian", lambda = 1) #LASSO
```

Random Forest

- Use the package `randomForest` or `ranger`.
 - `x`: input matrix.
 - `y`: target vector.
 - `ntree`: number of trees.
 - `mtry`: number of variables randomly sampled.
 - `replace`: sampling with or without replacement.
 - `cutoff`: threshold point (only for classification).
 - `nodesize`: maximum number of terminal nodes.

R Code

```
#Load package
library(randomForest)

data(iris)
set.seed(71)

iris.rf <- randomForest(
  x = as.matrix(iris[, c("Sepal.Length", "Sepal.Width", "Petal.Length", "Petal.Width")]),
  y = iris$Species, ntree = 500)
```

Neural Networks

- Use the package `nnet`.
 - `x`: input matrix.
 - `y`: target vector.
 - `linout`: switch for linear output units. Default logistic output units.
 - `entropy`: switch for entropy (= maximum conditional likelihood) fitting.
 - `softmax`: switch for softmax (log-linear model).
 - `censored`: A variant on softmax, in which non-zero targets mean possible classes. `linout`, `entropy`, `softmax` and `censored` are mutually exclusive.

R Code

```
#Load package
library(randomForest)

data(iris)
set.seed(71)

iris.nn <- nnet(
  x = as.matrix(iris[, c("Sepal.Length", "Sepal.Width", "Petal.Length", "Petal.Width")]),
  y = iris$Species, data=iris, maxit = 200)
```

Implementing ML Algorithms in R

- All the ML algorithms have easy-to-implement R packages/functions!

R Code

```
# LASSO
lasso <- glmnet(x = xtrain, y = ytrain, family = 'binomial', lambda = 0.1)
# check out the cv.glmnet also!

# random forest
rf <- randomForest(x = xtrain, y = factor(ytrain))

# neural network
nn <- nnet(x = xtrain, y = ytrain, size = 10, entropy = TRUE, trace = FALSE)

# support vector machine
library(e1071)

# svm linear kernel
svm_lin <- svm(x = xtrain, y = factor(ytrain), kernel = "linear")

# svm radial kernel
svm_rbf <- svm(x = xtrain, y = factor(ytrain), kernel = "radial")
```

Implementing ML Algorithms in Python

- All the ML algorithms have easy-to-implement Python libraries/functions!
- Popular Python libraries for machine learning: PyTorch, keras, scikit-learn, TensorFlow, etc.

Python Code

```
# LASSO

from glmnet import LogitNet
# Initialize and fit the LASSO model
lasso = LogitNet(alpha=1, lambda_path=np.array([0.1]))
lasso = lasso.fit(xtrain, ytrain)

# random forest

from sklearn.ensemble import RandomForestClassifier
# Initialize the RandomForestClassifier
rf = RandomForestClassifier()

# Fit the model to the data
rf.fit(xtrain, ytrain)
```

Implementing ML Algorithms in Python (Cont')

Python Code

```
# neural networks

from sklearn.neural_network import MLPClassifier

# Initialize the MLPClassifier
nn = MLPClassifier(hidden_layer_sizes=(10,), activation='relu', solver='adam',
                    max_iter=200, verbose=False, random_state=1)

# Fit the model to the data
nn.fit(xtrain, ytrain)

# support vector machine

from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

# SVM with Linear Kernel
svm_lin = make_pipeline(StandardScaler(), SVC(kernel='linear'))
svm_lin.fit(xtrain, ytrain)

# SVM with Radial Basis Function (RBF) Kernel
svm_rbf = make_pipeline(StandardScaler(), SVC(kernel='rbf')) # Note: 'rbf' is the default kernel
svm_rbf.fit(xtrain, ytrain)
```