

Causal Inference (6.S059/15.Co8/17.Co8)

Recitation, Week 9.

Topic: Causal Machine Learning

Licheng Liu

April 19, 2024

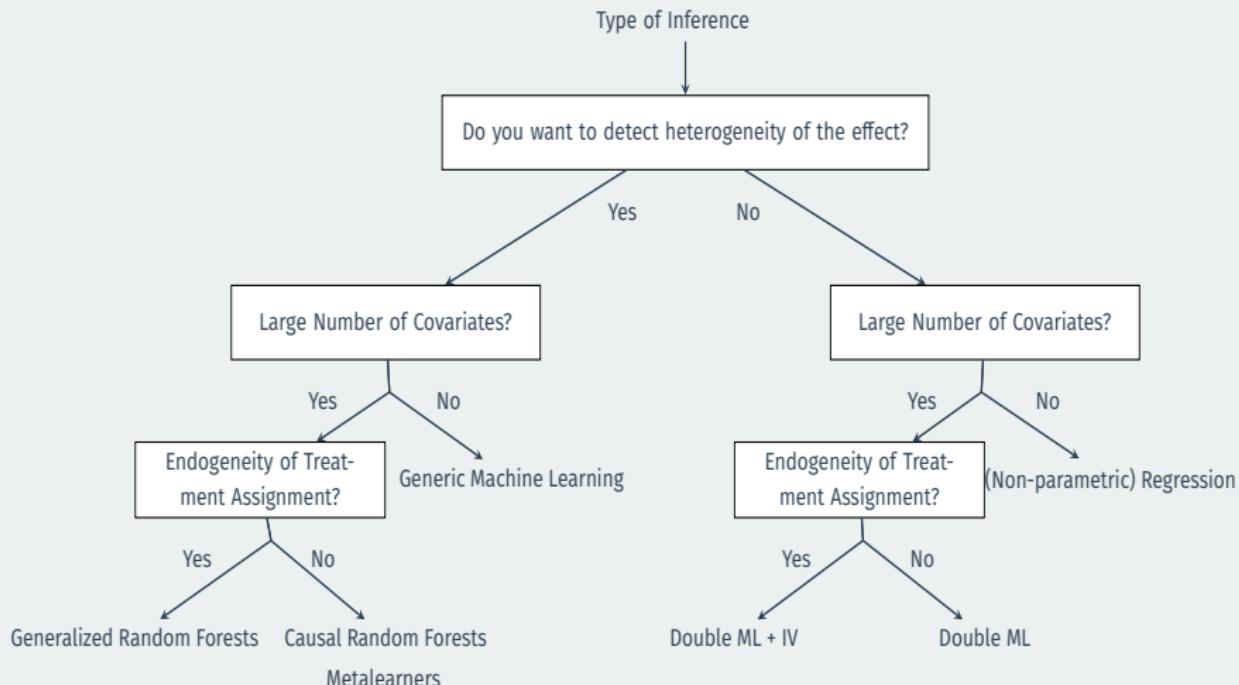
MIT

Table of contents

1. Theoretical Roadmap
2. High Dimensionality of X
3. Heterogeneous Treatment Effects

1/ Theoretical Roadmap

Causal Machine Learning



What Type of Problems Can Machine Learning Solve?

- Machine Learning techniques cannot solve causal identification problems.
 - How much can we learn about parameters from infinite amount of data? (Manski, 2007).
- The predictive power of ML tools can indeed solve some inferential issues.
 - How much can we learn about parameters from infinite amount of data? (Manski, 2007).
 - However, these techniques create potential new problems, which must be taken into consideration (efficiency, bias, etc.).

What kind of problems will we try to solve?

1. High Dimensionality of \mathbf{X} :

- I have a large number of covariates ($p >> n$). How can I use them to justify my identification strategy (Selection on Observables)?

2. Heterogeneous Treatment Effects:

- The typical Quantity of Interest is the Conditional Average Treatment Effect. But what if I want a different QoI? (Fundamental Problem of Causal Inference).
- A typical strategy is to use a linear interactive model. But these models impose strong assumptions (see Hainmueller, Mummolo, and Xu, 2019).

2/ High Dimensionality of X

High Dimensionality of X

- Definition: $p \gg n$.
- Typical Regression Setting: if $p > n$, then $\widehat{\beta}_{OLS}$ is not uniquely defined.
For $p < n$ but *large*, $\widehat{\beta}_{OLS}$ can be unstable and have high variance.
- ML based approach can:
 - deal with arbitrary interactions or flexible specification.
 - regularize irrelevant terms.
- **Sparsity:** only a subset of the elements of β are nonzero.
- **Approximate Sparsity:** generalization of Sparsity.
 - Recasts the High-Dimensional problem in a variable selection framework.

High Dimensionality of X

- Algorithms
 1. **Post-LASSO:** Choose correct variables for outcome model (naive-Lasso).
 2. **Double Selection:** Choose correct variables for treatment and outcome model.
 3. **Double/Debiased ML:** Choose correct variables for treatment and outcome model and allow for flexible functional form.
- Model Selection and Estimation cannot be achieved optimally at the same time.
 - Chances of mistakes in selection step = Contamination of estimation+inference step. Regularization Bias

Post Lasso (Naïve LASSO)

- One (naive) approach is to focus on just the outcome model.
- Instead of specifying a small set of covariates that almost all matter, analyst specifies a large number of covariates that mostly do not.
- Use a Lasso model to choose the ones that are significant (non-zero covariates).
- Coverage probability of standard confidence intervals can be far from the nominal level (downward bias). Affected by (and increasing in) the correlation between D and X Omitted-Variable Bias (Zero if $\beta = 0$ or $\rho = 0$).

Post Lasso (Naïve LASSO)

R Code

```
#Load package
library(glmnet)

# CV is not how chernozhukov et al choose

# Run the Outcome model
lasso <- cv.glmnet(x = X, y = y)

# Choose which variables to keep
keep <- as.matrix(coef(lasso, s = "lambda.min")[-1]) !=0

# Run OLS
lm(y ~ T + X[, keep])
```

Double Selection

- Concern with post Lasso: it might miss weak predictors of the outcome that are strong predictors of the treatment.
- This is particularly a concern if the outcome model is non-sparse or if covariates are highly correlated.
- Want to have the covariates and functional form suitable for outcome and treatment models.
- Solution is to select covariates that predict both the treatment and the outcome \leadsto Double Selection.
- You can put many covariates as you want as well as their interactions.
- Under approximate sparsity, asymptotic normality
- **Post-Regularization LASSO** (Partialling Out LASSO): Gains in efficiency. Uses only the relevant components of X to separately demean Y and D , leading to a greater parsimony.

Double Selection

R Code

```
# CV is not how chernozhukov et al choose
# the penalty term. decide outcome model
outcome_model <- cv.glmnet(x = X, y = y)

# Choose covariates whose coefficients are not zero
keep_outcome_model <- as.matrix(coef(outcome_model, s = "lambda.min")[-1]) != 0

# Decide treatment model
treat_model <- cv.glmnet(x = X, y = T, family = "binomial")

# Choose covariates whose coefficients are not zero
keep_treat_model <- as.matrix(coef(treat_model, s = "lambda.min")[-1]) != 0

# Run regression on chosen variables
lm(y ~ T + X[, keep_outcome_model|keep_treat_model])

### Alternative Code
library(hdm)

#### Implement Double Selection with LASSO
dsl <- rlassoEffect(x = X, d = T, y = Y, method = "double selection", post = TRUE)
summary(dsl)
```

Double ML

- In principle, double selection can be used to protect us from many forms of model misspecification.
- For example, a high order polynomial can fit most response functions very flexibly.
- However, no guarantees that this function is sparse or efficiently estimated.
- Double ML – tries to overcome this by using flexible ML methods to model the covariates.
- Suppose true model is:

$$Y = \tau T + g(X) + \epsilon$$

$$T = m(X) + \eta$$

Double ML

- If we knew $m(\cdot)$ and $g(\cdot)$, recovering τ would be trivial.
 \rightsquigarrow [FWL Theorem] Just regress $u = Y - g(X)$ on $e = T - m(X)$.
- Immunization/Orthogonalization Procedure.
- **Intuition:** remove a part of T and Y that can be explained by X (i.e., partialing out, obtain residual), and then run a regression of them.
- Instead, we will use machine learning to flexibly model \hat{g} and \hat{f} .
- The problem with this is overfitting – that we'll capture noise or the effect of the treatment in our estimates.
- The solution is **cross fitting** – fit $m(X)$ and $g(X)$ within one part of the sample and estimate the residuals on the other.
- Sample Splitting reduces the dependency between the estimation stages and can improve performance.

Double ML Implementation

R Code

```
# also check slide 11
# for each fold k,
#get_resids <- function(X,y,treat, fold, folds) {
#    d <- data.frame(y = y[fold != folds], X = X[fold != folds,])
#    outcome_model <- ranger(y ~ ., data = d)
#    d <- data.frame(treat = treat[fold != folds], X = X[fold != folds,])
#    treat_model <- ranger(treat ~ ., data = d)
#    V_hat <- treat[fold == folds] - predict(treat_model, newx = X[fold == folds,])
#    W_hat <- y[fold == folds] - predict(outcome_model, newx=X[fold == folds,])
#    mod <- lm(W_hat~V_hat)
#    tau <- coef(mod)[2]
#    epsilon <- resid(mod)
#    return(list(tau = tau, epsilon = epsilon, v_hat = V_hat))
#}
# naive approach to conduct k-fold cross validation
folds <- sample(1:k, nrow(X), replace = TRUE)
Vsqrdf <- 0 VTimesEpsilon <- 0 tau <- rep(NA, k) for(i in 1:k) {
    results <- get_resids(X, y, treat = treat, fold = i, folds = folds)
    tau[i] <- results$tau
    Vsqrdf <- Vsqrdf + sum(results$v_hat^2)
    VTimesEpsilon <- VTimesEpsilon + t(results$v_hat^2) %*% results$epsilon^2
}
tau <- mean(tau)
sesqrdf <- (Vsqrdf/nrow(X))^{(-2)}*(VTimesEpsilon/nrow(X))
se <- sqrt(sesqrdf/nrow(X))
```

Double ML Implementation

R Code

```
### Alternative Code
library(DoubleML)
library(mlr3)
library(mlr3learners)

#### Define data for DDML
dml_data_sim <- double_ml_data_from_matrix(X = X, y= Y, d = T)

#### Select LASSO learner
learner <- lrn("regr.cv_glmnet", s="lambda.min")
ml_l_sim <- learner$clone()
ml_m_sim <- learner$clone()

#### Execute DDML
obj_dml_plr_sim <- DoubleMLPLR$new(dml_data_sim, ml_l=ml_l_sim, ml_m=ml_m_sim)
obj_dml_plr_sim$fit()
print(obj_dml_plr_sim)
```

3/ Heterogeneous Treatment Effects

Conditional Average Treatment Effects

- CATE: $\tau(D; X) = \mathbb{E}(Y_i^1 - Y_i^0 | X_i = x)$

1. Linear Interactive Model

$$Y = \beta_0 + \beta_1 D + \beta_2 M + \beta_3 D \times M + \gamma X + \epsilon$$

- **Extra Assumptions:** the treatment effect varies linearly with M, and for each value of M, the treatment and control groups should have a sufficient number of overlapping cases (see **Hainmueller, Mummolo, and Xu, 2019**).
2. Machine learning tools to estimate the heterogeneous treatment effects:
- Loss function approach: Squared Loss SVM with separate LASSO constraints (**Imai and Ratkovic, 2013**), R-learner (**Nie and Wager, 2021**).
 - Construct potential outcomes: X-learner (**Künzel et al., 2019**), Causal Forests (**Wager and Athey, 2018**).

Causal Forests

- **Tree:** method that recursively partitions the high-dimensional covariate space into smaller units.
- Prediction \neq Detection of Heterogeneity.
 - Minimization of within-leaf variance of Y (all the cases belonging to the same leaf should be homogeneous in terms of the outcome) \leadsto Minimization of the within-leaf variance of the estimated treatment effects (inter-leaf variation should be large).
 - New Splitting Rule:

$$\widehat{-\text{EMSE}_\tau}(S^{\text{Tr}}, N^{\text{est}}, \Pi) = \underbrace{\frac{1}{N^{\text{Tr}}} \sum_{i \in S^{\text{Tr}}} \hat{t}^2(X_i | S^{\text{Tr}}, \Pi)}_{\substack{\text{Variance of Treatment Effects} \\ \text{across Leaves} \\ \text{Prefer leaves with} \\ \text{Heterogeneous Effects}}} - \underbrace{\left(\frac{1}{N^{\text{Tr}}} + \frac{1}{N^{\text{est}}} \right) \sum_{I \in \Pi} \left(\frac{S_{S^{\text{Tr}}_{\text{Treat}}}^2(I)}{p} + \frac{S_{S^{\text{Tr}}_{\text{Control}}}^2(I)}{1-p} \right)}_{\substack{\text{Uncertainty about Leaf} \\ \text{Treatment Effects} \\ \text{Prefer leaves with Good Fit} \\ (\text{Leaf-Specific Effects estimated} \\ \text{Precisely})}}$$

Causal Forests

- **Honest Procedure:** not use the same information for selecting the model structure as for estimation given a model structure.
 - One (independent) split of the data is used to learn the tree structure/-partition, and the second split of the data is used to conduct inference (estimation of treatment effects).

R Code

```
# Run Causal Forests (Basic Algorithm)
library(grf)

# Estimate Causal Forest
cf <- causal_forest(X = X, Y = Y, W = W, num.trees = 10000,
                      honesty = TRUE, honesty.fraction = 0.5,
                      tune.parameters = "all", seed = 17806)

# Estimate Predicted Values (CATEs)
pred <- predict(object = cf, newdata = newX,
                  estimate.variance = TRUE)$predictions

# Use expanded algorithm for Problem 3.
```

Metalearners

- Meta-algorithms decompose estimating the CATE into several subregression problems that can be solved with any supervised ML method.
 - Combination of **base learners** in a specific manner while allowing the base learners to take any form.
1. **S-Learner** (Single) \leadsto using all of the features and the treatment indicator (without giving to D a special role).

$$\mu(x) = \mathbb{E}[Y|X = x, D = d]$$

$$\hat{\tau}(x) = \hat{\mu}(x, D = 1) - \hat{\mu}(x, D = 0)$$

- Risk of dropping the treatment.
- Low statistical efficiency.

Metalearners

2. **T-Learner** (Two) \leadsto use base learners to estimate the conditional expectations of the outcomes separately for control and treatment groups.

$$\mu_0(x) = \mathbb{E}[Y|D=0, X=x]$$

$$\mu_1(x) = \mathbb{E}[Y|D=1, X=x]$$

$$\hat{\tau}(x) = \hat{\mu}_1(x) - \hat{\mu}_0(x)$$

- Ignore group size.
- Prediction \neq Heterogeneity.

Metalearners

3. **X-Learner** ↗ uses each observation in the training set in an X-like shape (Sample Splitting for Fundamental Problem of Causal Inference).

3.1 Estimate the response functions:

$$\mu_0(x) = \mathbb{E}[Y|D = 0, X = x]$$

$$\mu_1(x) = \mathbb{E}[Y|D = 1, X = x]$$

- 3.2 Impute the individual treatment effects (for the Treated group with the control-outcome estimator and for the Control group with the treatment-outcome estimator):

$$\tilde{D}_i^1 = Y_i^1 - \hat{\mu}_0(X_i^1)$$

$$\tilde{D}_i^0 = \hat{\mu}_1(X_i^0) - Y_i^0$$

Metalearners

3.3 Use any learner to estimate/predict the imputed treatment effects for each group:

$$\hat{\tau}_1(x) = \mathbb{E}[\tilde{D}_i^1 | D = 1, X = x]$$

$$\hat{\tau}_0(x) = \mathbb{E}[\tilde{D}_i^0 | D = 0, X = x]$$

3.4 Define the CATE estimate by the weighted average of the two estimates in 3.3:

$$\hat{\tau}(x) = g(x)\hat{\tau}_0(x) + (1 - g(x))\hat{\tau}_1(x)$$

- $g(x) \in [0, 1]$. Good option =Propensity Score.
- Efficient use of unbalanced design.